

Introduction to Supervised Learning

From Features to Targets

Mehmet Mercangöz

AI and AIS Summer School July/Aug 2026

① Foundations

- supervised learning, labels, features, regression, classification

② Learning as optimisation

- loss functions, least squares, polynomial regression

③ Generalisation and evaluation

- train/validation/test split, overfitting, metrics

④ Classification models

- logistic regression, decision thresholds, confusion matrix

⑤ Practical modelling choices

- data preparation, model selection

What Is Supervised Learning?

Supervised learning uses examples where the desired answer is already known.

Each training example contains:

input features $\mathbf{x}_i$ $\longrightarrow$ target or label y_i .

The goal is to learn a prediction rule

$$\hat{y} = f(\mathbf{x})$$

that gives useful predictions for **new, unseen** inputs.

Core idea

The model should not merely remember the training data. It should **generalise**.

Supervised, Unsupervised, and Reinforcement Learning

Setting	Data	Typical question
Supervised learning	Inputs with known targets	Can we predict the target for a new input?
Unsupervised learning	Inputs without targets	Are there patterns, groups, or structure?
Reinforcement learning	Actions, states, rewards	Which actions should be taken over time?

This lecture focuses on **supervised learning**.

Regression and Classification

Supervised learning problems usually fall into two main types.

Regression

Predict a numerical value.

Examples:

- house price;
- energy demand;
- product quality;
- reactor conversion.

Classification

Predict a category or class.

Examples:

- spam or not spam;
- faulty or healthy;
- image class;
- high-risk or low-risk.

The Supervised Learning Dataset

Suppose we have n examples and p features.

$$X = \begin{bmatrix} - & - & - & \mathbf{x}_1^\top & - & - & - \\ - & - & - & \mathbf{x}_2^\top & - & - & - \\ & & & \vdots & & & \\ - & - & - & \mathbf{x}_n^\top & - & - & - \end{bmatrix} = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1p} \\ x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \cdots & x_{np} \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}.$$

- Each **row** is one observation.
- Each **column** is one feature.
- The target y_i is what we want to predict.

A Small Example Dataset

Observation	Temperature	Pressure	Quality
1	80	2.1	91.2
2	83	2.4	92.8
3	77	1.9	89.7
4	85	2.5	94.1

Here, a possible supervised learning task is:

$(\text{temperature}, \text{pressure}) \mapsto \text{quality}.$

Question

Is this a regression problem or a classification problem?

The Supervised Learning Workflow

- 1 Define the prediction problem.
- 2 Identify the inputs $\mathbf{x}$ and target y .
- 3 Prepare and inspect the data.
- 4 Split the data into training, validation, and test sets.
- 5 Choose a model class.
- 6 Train the model by minimising a loss function.
- 7 Evaluate the trained model on unseen data.
- 8 Inspect errors and improve the model if needed.

Important

The algorithm is only one part of the workflow. Problem formulation and data quality are often just as important.

What Does It Mean for a Model to Learn?

A model usually contains adjustable parameters.

$$\hat{y}_i = f(\mathbf{x}_i; \boldsymbol{\theta}).$$

Training means choosing parameters $\boldsymbol{\theta}$ so that predictions match the observed targets as well as possible.

$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta}} L(\boldsymbol{\theta}).$$

- $L(\boldsymbol{\theta})$ is the **loss function**.
- It measures prediction error on the training data.
- Supervised learning is therefore closely connected to optimisation.

Identify: Problem Type, Targets, Features

Application

Regression or Classification?

Predict tomorrow's electricity demand using weather and calendar data

Determine whether an incoming email is spam

Estimate product purity from process measurements

Determine whether a machine is healthy or faulty

Linear Regression

Linear regression predicts a numerical target using a weighted sum of features.

$$\hat{y}_i = w_0 + w_1x_{i1} + w_2x_{i2} + \cdots + w_px_{ip}.$$

Using a bias column of ones, this can be written compactly as

$$\hat{\mathbf{y}} = X\mathbf{w}.$$

- X contains the input data.
- $\mathbf{w}$ contains the coefficients to be estimated.
- $\hat{\mathbf{y}}$ contains the model predictions.

Residuals and Least-Squares Loss

For each training example, the residual is

$$e_i = y_i - \hat{y}_i.$$

For all examples together:

$$\mathbf{e} = \mathbf{y} - X\mathbf{w}.$$

The least-squares objective is

$$J(\mathbf{w}) = \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \|\mathbf{y} - X\mathbf{w}\|_2^2.$$

Choose the coefficients that make the total squared prediction error as small as possible.

Least-Squares Solution

The least-squares problem is

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \|\mathbf{y} - X\mathbf{w}\|_2^2.$$

Expanding the objective:

$$J(\mathbf{w}) = (\mathbf{y} - X\mathbf{w})^\top (\mathbf{y} - X\mathbf{w}).$$

Taking the gradient and setting it equal to zero gives

$$\nabla_{\mathbf{w}} J = 2X^\top (X\mathbf{w} - \mathbf{y}) = \mathbf{0}.$$

Therefore,

$$X^\top X\mathbf{w} = X^\top \mathbf{y}.$$

These are called the **normal equations**.

Closed-Form Linear Regression Estimate

If $X^\top X$ is invertible, the least-squares estimate is

$$\hat{\mathbf{w}} = (X^\top X)^{-1} X^\top \mathbf{y}$$

This is a direct solution of an optimisation problem.

When is this formula safe to use?

It requires $X^\top X$ to be invertible and reasonably well-conditioned.

If this is not the case, we often use the pseudoinverse:

$$\hat{\mathbf{w}} = X^+ \mathbf{y}.$$

Least Squares: Geometric Interpretation

The model predictions have the form

$$\hat{\mathbf{y}} = X\hat{\mathbf{w}}.$$

This means $\hat{\mathbf{y}}$ must lie in the column space of X .

Least squares chooses the point in that space that is closest to $\mathbf{y}$.

$$\mathbf{y} = \hat{\mathbf{y}} + \mathbf{e}, \quad X^T \mathbf{e} = \mathbf{0}.$$

Key idea

The residual vector is orthogonal to every column of the design matrix.

Linear Regression: Parameters and Hyperparameters

In linear regression, the coefficients are **parameters**.

$$\mathbf{w} = [w_0 \quad w_1 \quad \cdots \quad w_p]^\top.$$

They are learned from the data by solving an optimisation problem.

A **hyperparameter** is chosen by the modeller before training.

Examples:

- the degree of a polynomial regression model;
- the number of neighbours in k -nearest neighbours;
- the maximum depth of a decision tree;
- the strength of regularisation.

Linear Regression: An Exercise



`https://aisl-at-imperial-college-london.github.io/summer\_school\_repo/man\_opt\_LR.html`

Polynomial Regression

Suppose the relationship between x and y is curved.

A polynomial regression model may use

$$\hat{y} = a_0 + a_1x + a_2x^2 + a_3x^3.$$

This is nonlinear as a function of x , but it is still **linear in the unknown coefficients**:

$$\mathbf{a} = [a_0 \quad a_1 \quad a_2 \quad a_3]^\top.$$

Therefore, estimating the coefficients is still a linear least-squares problem.

Polynomial Regression as Linear Least Squares

For training data (x_i, y_i) , define a new design matrix:

$$\Phi = \begin{bmatrix} 1 & x_1 & x_1^2 & x_1^3 \\ 1 & x_2 & x_2^2 & x_2^3 \\ \vdots & \vdots & \vdots & \vdots \\ 1 & x_n & x_n^2 & x_n^3 \end{bmatrix}.$$

The model becomes

$$\hat{\mathbf{y}} = \Phi \mathbf{a}.$$

The coefficient estimate is

$$\hat{\mathbf{a}} = (\Phi^\top \Phi)^{-1} \Phi^\top \mathbf{y}$$

when $\Phi^\top \Phi$ is invertible.

Polynomial Regression: Important Lesson

Nonlinear features do not necessarily mean nonlinear optimisation.

A model can be nonlinear in the input x , but linear in the unknown coefficients.

Examples of linear-in-parameters models:

$$\hat{y} = a_0 + a_1x + a_2x^2,$$

$$\hat{y} = a_0 + a_1 \sin(x) + a_2 \cos(x),$$

$$\hat{y} = a_0 + a_1T + a_2P + a_3TP.$$

All can be fitted using least squares after constructing suitable features.

Model Complexity and Polynomial Degree

The polynomial degree is a modelling choice.

Degree	Model flexibility	Risk
1	Low	Underfitting
2–4	Moderate	Often useful
Very high	High	Overfitting

Key idea

A more flexible model may fit the training data better, but that does not guarantee better predictions on new data.

Optional Extension: Regularised Least Squares

If a model is too flexible or the data matrix is ill-conditioned, we can penalise large coefficients.

$$J(\mathbf{w}) = \|\mathbf{y} - X\mathbf{w}\|_2^2 + \lambda\|\mathbf{w}\|_2^2.$$

The solution is

$$\hat{\mathbf{w}} = (X^\top X + \lambda I)^{-1} X^\top \mathbf{y}$$

- $\lambda = 0$: ordinary least squares.
- Larger λ : stronger penalty on large coefficients.

Training, Validation, and Test Sets

We usually split the data into three parts.

Dataset part	Purpose
Training set	Used to fit the model parameters.
Validation set	Used to compare models and tune hyperparameters.
Test set	Used only at the end to estimate performance on unseen data.

Warning

Repeatedly tuning a model based on test-set performance makes the test set no longer independent.

Generalisation

A model generalises well if it performs well on new data drawn from the same underlying problem.

Case	Description
Underfitting	Model too simple; poor training and test performance.
Good fit	Captures the main pattern; good test performance.
Overfitting	Model follows noise; excellent training performance but poor test performance.

The objective is not to minimise training error at all costs. The objective is useful prediction on unseen data.

Regression Metrics

For regression, common error metrics include:

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|,$$

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2,$$

$$\text{RMSE} = \sqrt{\text{MSE}}.$$

- MAE is easy to interpret and less sensitive to large errors.
- MSE penalises large errors more strongly.
- RMSE has the same units as the target variable.

The Coefficient of Determination: R^2

The coefficient of determination compares the model to a simple baseline that always predicts the mean.

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}.$$

- $R^2 = 1$: perfect predictions.
- $R^2 = 0$: no better than predicting the mean.
- $R^2 < 0$: worse than predicting the mean.

Caution

A high R^2 on the training data does not necessarily imply good generalisation.

Classification: Scores, Probabilities, and Classes

In binary classification, the target is often

$$y_i \in \{0, 1\}.$$

Many classifiers first compute a score or probability:

$$p_i = P(y_i = 1 \mid \mathbf{x}_i).$$

Then a decision threshold is applied:

$$\hat{y}_i = \begin{cases} 1, & p_i \geq 0.5, \\ 0, & p_i < 0.5. \end{cases}$$

The threshold can be changed when false positives and false negatives have different costs.

Logistic Regression

Logistic regression is a supervised learning model for classification.

First compute a linear score:

$$z_i = \mathbf{w}^\top \mathbf{x}_i + b.$$

Then convert the score into a probability using the sigmoid function:

$$p_i = \sigma(z_i) = \frac{1}{1 + e^{-z_i}}.$$

Interpretation

p_i is the model's estimated probability that example i belongs to class 1.

Why Not Use Linear Regression for Classification?

Linear regression predicts

$$\hat{y} = \mathbf{w}^\top \mathbf{x} + b.$$

For classification, this can be problematic:

- predictions can be below 0 or above 1;
- squared error is not naturally matched to binary labels;
- we often want calibrated probabilities, not just numerical scores.

Logistic regression instead uses

$$0 < \sigma(z) < 1,$$

so the model output can be interpreted as a probability.

Logistic Regression Loss Function

For binary labels $y_i \in \{0, 1\}$, logistic regression uses the negative log-likelihood:

$$J(\mathbf{w}, b) = - \sum_{i=1}^n [y_i \log(p_i) + (1 - y_i) \log(1 - p_i)] .$$

where

$$p_i = \sigma(\mathbf{w}^\top \mathbf{x}_i + b).$$

- If $y_i = 1$, the loss is small when p_i is close to 1.
- If $y_i = 0$, the loss is small when p_i is close to 0.

Logistic Regression as an Optimisation Problem

Logistic regression estimates its parameters by solving

$$(\mathbf{w}^*, b^*) = \arg \min_{\mathbf{w}, b} J(\mathbf{w}, b).$$

Unlike ordinary least squares, there is generally no closed-form solution like

$$(X^\top X)^{-1} X^\top \mathbf{y}.$$

Instead, we use iterative numerical optimisation.

Important contrast

Linear regression has a closed-form least-squares solution. Logistic regression is usually fitted by an optimisation algorithm.

Logistic Regression: Gradient Form

Using augmented feature vectors

$$\tilde{\mathbf{x}}_i = \begin{bmatrix} 1 \\ \mathbf{x}_i \end{bmatrix},$$

the predicted probability for observation i is

$$p_i = \sigma\left(\boldsymbol{\theta}^\top \tilde{\mathbf{x}}_i\right).$$

Collecting all predictions,

$$\mathbf{p} = [p_1 \quad p_2 \quad \cdots \quad p_n]^\top.$$

The gradient of the average logistic loss is

$$\nabla_{\boldsymbol{\theta}} J(\boldsymbol{\theta}) = \frac{1}{n} \tilde{X}^\top (\mathbf{p} - \mathbf{y}).$$

Therefore, gradient descent gives

$$\boldsymbol{\theta}^{(k+1)} = \boldsymbol{\theta}^{(k)} - \frac{\alpha}{n} \tilde{X}^\top (\mathbf{p}^{(k)} - \mathbf{y})$$

Decision Boundaries

In logistic regression, the class prediction changes when

$$p_i = 0.5.$$

Since

$$\sigma(z_i) = 0.5 \iff z_i = 0,$$

we get the decision boundary

$$\mathbf{w}^\top \mathbf{x} + b = 0.$$

- In two dimensions, this is a line.
- In three dimensions, this is a plane.
- In higher dimensions, this is a hyperplane.

Confusion Matrix

For binary classification, predictions can be summarised as:

	Predicted positive	Predicted negative
Actually positive	True positive	False negative
Actually negative	False positive	True negative

- A false positive is a false alarm.
- A false negative is a missed detection.
- Which one is worse depends on the application.

Classification Metrics

From the confusion matrix, we can compute:

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN},$$

$$\text{Precision} = \frac{TP}{TP + FP},$$

$$\text{Recall} = \frac{TP}{TP + FN},$$

$$F_1 = 2 \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}}.$$

Accuracy can be misleading when one class is much more common than the other.

Threshold Choice Matters

The default threshold for binary classification is often 0.5:

$$\hat{y} = 1 \quad \text{if} \quad p \geq 0.5.$$

But this threshold is a decision choice, not a law of nature.

Lower threshold	More positives predicted; usually higher recall, more false positives.
Higher threshold	Fewer positives predicted; usually higher precision, more false negatives.

In safety-critical fault detection, a missed fault may be more costly than a false alarm.

Receiver Operating Characteristic (ROC) Curve

Changing the classification threshold changes the balance between **detected positives** and **false alarms**.


$$\text{TPR} = \frac{\text{TP}}{\text{TP} + \text{FN}}, \quad \text{FPR} = \frac{\text{FP}}{\text{FP} + \text{TN}}.$$

The ROC curve plots

TPR against FPR

over all possible thresholds.

- Upper-left: ideal performance
- Diagonal: random guessing
- Larger AUC: better discrimination



AUC = area under the ROC curve

0.5 : random, 1 : perfect.

Data Preparation

Before training a supervised learning model, we often need to handle:

- missing values;
- categorical variables;
- different feature scales;
- outliers;
- duplicated or inconsistent records;
- imbalanced classes;
- data leakage.

Data leakage

Data leakage occurs when information is used during training that would not be available at prediction time.

Feature Scaling

Some models are sensitive to the numerical scale of features.

For example, k -nearest neighbours uses distances:

$$\|\mathbf{x}_i - \mathbf{x}_j\|_2.$$

If one feature is measured in thousands and another is measured between 0 and 1, the large-scale feature may dominate the distance.

A common standardisation is

$$x_{ij}^{\text{scaled}} = \frac{x_{ij} - \mu_j}{s_j},$$

where μ_j and s_j are computed from the training data.

Model Selection

Different models and hyperparameters can be compared using validation data.

Choice	Examples
Model class	linear regression, tree, logistic regression, k -NN
Hyperparameters	polynomial degree, tree depth, number of neighbours
Metric	MAE, RMSE, accuracy, precision, recall

Rule

Use the validation set for model selection. Use the test set only for final assessment.

Use of Supervised Learning with Care!

A supervised learning model can be wrong for many reasons.

- The training data may be biased or incomplete.
- The future may differ from the past.
- The selected features may miss important information.
- The model may overfit noise.
- Different types of error may have very different consequences.

Practical message

Always ask: what happens if the model is wrong?

End-to-End Worked Example Structure

A complete supervised learning example can follow this structure.

- 1 Define the prediction task.
- 2 Identify features and target.
- 3 Inspect and prepare the data.
- 4 Split into training, validation, and test sets.
- 5 Train a simple baseline model.
- 6 Evaluate on validation data.
- 7 Try a more flexible model.
- 8 Compare metrics and inspect errors.
- 9 Evaluate once on the test set.
- 10 Decide whether the model is useful for the intended purpose.

Summary

- Supervised learning learns from labelled examples.
- Regression predicts numerical values; classification predicts categories.
- Training usually means minimising a loss function.
- Linear regression can be solved by least squares.
- Polynomial regression is linear in the coefficients after feature transformation.
- Logistic regression is fitted by iterative numerical optimisation.
- Generalisation matters more than training error.
- Evaluation metrics must match the purpose of the model.